

From Problem to Coroutine: Reducing I/O Latency

CHEINAN MARKS



About This Talk

- Francesco Zoffoli: <https://github.com/MakersF/cppcon-2021-corobatch>
- 2016-2019 CppCon Talks by Gor Nishanov and James McNellis
- Lewis Baker's Blogs: <https://lewissbaker.github.io/>
- N4868 C++ 20 Draft:
<https://github.com/cplusplus/draft/releases/download/n4868/n4867.html>
- cheinan@pobox.com cyliber.org



I Had a Simple Program to Write

- Gather DNA Reads from FASTQ
- Count Instances of kmers



I Had a Simple Program to Write

- Gather DNA reads from **FASTQ**
- Count instances of kmers
- Simple(?) Line-Based ASCII File in 4 line records
- DNA Read is second line
- Read is ASCII string made of A, C, G, T, N



I Had a Simple Program to Write

- Gather DNA reads from FASTQ
- Count instances of **kmers**
- Polymer – multiple repetition of single molecule
- E.g. polyethylene, our chromosomes
- Monomer – polymer building block
- Dimer – pair of monomers
- Kmer – k (arbitrary integer) number of monomers
- E.g. AATGGGC for $k = 7$



I Had a Simple Program to Write

- Gather DNA Reads from FASTQ
- Count Instances of kmers
- Substring search
- Count instances of “AT” in “GATGATC”



I Had a Simple Program to Write

- Gather DNA Reads from FASTQ
- Count Instances of kmers
- Substring search
- Count instances of “AT” in “GATGATC”
- Answer: 2



I Had a Simple Program to Write

```
while(stream.ignore(maxSize, '\n')) {  
    std::getline(stream, line);  
    kmerCount += CountKmer(line, kmer);  
    stream.ignore(maxSize, '\n');  
    stream.ignore(maxSize, '\n');  
}  
  
return kmerCount;
```




I Had a Simple Program to Write

```
std::string_view::size_type kmerPos = 0;
while ((
    kmerPos =
        line.find(kmer, kmerPos)) != line.npos) {
    kmerCount++;
    kmerPos++;
}
return kmerCount;
```



I Had a Simple Program to Write

- Gather DNA Reads from FASTQ
 - Read lines from a file
 - Extract a portion of a record
 - I/O
- Count Instances of kmers
 - Compute on the record



I Had a Simple Program to Write

- Gather DNA Reads from FASTQ
 - Read lines from a file
 - Extract a portion of a record
 - I/O
- Count Instances of kmers
 - Compute on the extracted data
- Producer – Consumer threads



I Had a Simple Program to Write

- Gather DNA Reads from FASTQ
 - Read lines from a file
 - Extract a portion of a record
 - I/O
- Count Instances of kmers
 - Compute on the extracted data
- Producer – Consumer threads
- Zzzzzzz



I Had a Great Idea

- Use coroutines to reduce latency



I Had a Great Idea

- Use coroutines to reduce latency
- It's easy, if convoluted
 - Add `co_await` keyword to async parts
 - Add a bunch of boilerplate code
- Enjoy Speedup
 - And get practice with a new C++20 feature



Expectation

```
std::string read = GetRead(stream);  
int kmerCount = CountKmers(read);
```



Expectation

```
std::string read = co_await GetRead(stream);  
int kmerCount = co_await CountKmers(read);
```




Expectation

```
std::string read = co_await GetRead(stream);  
int kmerCount = co_await CountKmers(read);  
  
// Everything magically runs faster!
```



Learnings

- Consider not using coroutines
- Wait for standard library support in C++23 or C++26
 - Consider using a coroutine library if you can
- Plan out the use of the `co_*` keywords first in coro body
 - Then you can add the boilerplate
- The work goes in the coroutine body
- Be very careful with object lifetime and ownership
- You will have to write some sort of manager/queue
- Leave optimization for the very end



Consider Not Using Coroutines

The facilities the C++ Coroutines TS provides in the language can be thought of as a low-level assembly-language for coroutines. These facilities can be difficult to use directly in a safe way and are mainly intended to be used by library-writers to build higher-level abstractions that application developers can work with safely.

– Lewis Baker

<https://lewissbaker.github.io/2017/11/17/understanding-operator-co-await>



Learnings

- Consider not using coroutines
- Wait for standard library support in C++23 or C++26
 - Consider using a coroutine library if you can
- Plan out the use of the `co_*` keywords first in coro body
 - Then you can add the boilerplate
- The work goes in the coroutine body
- Be very careful with object lifetime and ownership
- You will have to write some sort of manager/queue
- Leave optimization for the very end



Learnings

- Consider not using coroutines
- Wait for standard library support in C++23 or C++26
 - Consider using a coroutine library if you can
- Plan out the use of the `co_*` keywords first in coro body
 - Then you can add the boilerplate
- The work goes in the coroutine body
- Be very careful with object lifetime and ownership
- You will have to write some sort of manager/queue
- Leave optimization for the very end



Plan out the use of the `co_*` keywords first

- What am I doing?
- Where can I use coroutines?
 - Should I use coroutines?



Plan out the use of the `co_*` keywords first

- What am I doing?
- Where can I use coroutines?
 - Should I use coroutines?
- Reading I/O in unstructured records
 - Can't calculate record location
- Compute on data received from I/O



Plan out the use of the `co_*` keywords first

- What am I doing?
- Where can I use coroutines?
 - What do the coroutines do?
 - Should I use coroutines?
- Reading I/O in unstructured records
 - Can't calculate record location
- Compute on data received from I/O



Where Can I Use Coroutines

- Offload compute to a worker thread



Where Can I Use Coroutines

- Offload compute to a worker thread

```
CountKmersTask GetReadAndCount(...) {  
    auto read = GetRead();  
    co_await ResumeInComputePool();  
    kmerCount = CountKmersInRead(read, kmer);  
    co_return kmerCount;  
}
```



Where Can I Use Coroutines

```
CountKmersTask GetReadAndCount(...) {  
    auto read = GetRead();  
    co_await ResumeInComputePool();  
    kmerCount = CountKmersInRead(read, kmer);  
  
    co_return kmerCount;  
}
```

- Transfer



Transfer

- Coroutines are *not* functions (subroutines)
- They exist independently of a thread or stack
- Can start on one thread and end on another
- No synchronization needed
- Can even transfer to another coroutine
 - State machines

```
co_await ResumeInComputePool();
```



Transfer

```
struct ResumeInComputePool {  
    bool await_ready() { return false; }  
    void await_suspend(std::coroutine_handle<> h)  
    {  
        pool.execute([](auto h) {h.resume();}, h);  
    }  
    void await_resume() {}  
};
```



Plan out the use of the `co_*` keywords first

- What am I doing?
- Where can I use coroutines?
 - Should I use coroutines?
- Reading I/O in unstructured records
 - Can't calculate record location
- Compute on data received from I/O



Where Can I Use Coroutines?

- Transform unstructured file records into DNA reads
- Extract data from file records
- Generator



Generator

```
FastqGenerator GetARead(  
    const std::filesystem::path &fastqPath) {  
    std::ifstream fastqStream(fastqPath);  
    std::string read;  
  
    while (fastqStream) {  
        fastqStream.ignore(...);  
        std::getline(fastqStream, read);  
        co_yield std::move(read);  
        fastqStream.ignore(...);  
        fastqStream.ignore(...);  
    }  
}
```




Generator Boilerplate

```
struct FastqGen {  
    FastqGen(const FastqGen &) = delete;  
    FastqGen &operator=(const FastqGen &) = delete;  
  
    FastqGen(FastqGen &&rhs) noexcept : h_(rhs.h_) { rhs.h_ = {}; }  
    FastqGen &operator=(FastqGen &&rhs) noexcept {  
        if (this != &rhs) {  
            if (h_) {  
                h_.destroy();  
            }  
            h_ = rhs.h_;  
            rhs.h_ = {};  
        }  
        return *this;  
    }  
}
```



Generator Boilerplate

```
struct FastqGenerator {  
    ~FastqGenerator() {  
        if (h_) {  
            h_.destroy();  
        }  
    }  
private:  
    FastqGenerator(handle h) : h_(h) {}  
    handle h_;  
};
```



Generator Boilerplate

```
struct FastqGenerator {  
    struct promise_type {  
        auto yield_value(std::string read) {  
            read_ = std::move(read);  
            return std::suspend_always{};  
        }  
  
        std::string read_;  
    };  
};
```



Generator Boilerplate

```
struct FastqGenerator {  
    bool move_next() {  
        if (h_ && !h_.done()) {  
            h_.resume();  
            return !h_.done();  
        }  
        return false;  
    }  
    std::string current_value() {  
        return std::move(h_.promise().read_);  
    }  
};
```



Generator Boilerplate Bonus

```
class FastqIter {
public:
    FastqIter &operator++() {
        coroutine_.resume();
        return *this;
    }

    const std::string &operator*() const { return coroutine_.promise().read_; }

    bool operator==(std::default_sentinel_t) const {
        return !coroutine_ || coroutine_.done();
    }

    explicit FastqIter(const handle coroutine) : coroutine_{coroutine} {}

private:
    handle coroutine_;
};

FastqIter begin() {
    if (h_) {
        h_.resume();
    }
    return FastqIter{h_};
}

std::default_sentinel_t end() { return {}; }
```



Learnings

- Consider not using coroutines
- Wait for standard library support in C++23 or C++26
 - Consider using a coroutine library if you can
- Plan out the use of the `co_*` keywords first in coro body
 - Then you can add the boilerplate
- **The work goes in the coroutine body**
- Be very careful with object lifetime and ownership
- You will have to write some sort of manager/queue
- Leave optimization for the very end



The work goes in the coroutine body

- A coroutine returns a class called a task or a job
- Truly async work (that notifies) can go in awaiter
 - Work in `await_suspend`
 - Notification calls `resume` (possibly on other thread)
- Needed sync work goes in coroutine



The work goes in the coroutine body

```
CountKmersTask GetReadAndCount(...) {  
    auto kmerCount = 0;  
  
    if (!isCancelled && readGenerator.move_next()) {  
        read = readGenerator.current_value();  
        if (read.empty()) {  
            co_return 0;  
        }  
        ... // Compute transfer and work  
    }  
}
```




The work goes in the coroutine body

```
CountKmersTask GetReadAndCount(...) {  
    std::string read;  
    // co_await ResumeInIOThread{};  
    if (!isCancelled && readGenerator.move_next()) {  
        read = readGenerator.current_value();  
        if (read.empty()) {  
            co_return 0;  
        }  
        ... // Compute transfer and work  
    }  
}
```



The work goes in the coroutine body

```
CountKmersTask GetReadAndCount(..., read) {  
    co_await ResumeInComputeThread{};  
    ...  
}
```



The work goes in the coroutine body

```
CountKmersTask GetReadAndCount(...) {  
    std::string read;  
    // co_await ResumeInIOThread{};  
    ...  
}
```



Learnings

- Consider not using coroutines
- Wait for standard library support in C++23 or C++26
 - Consider using a coroutine library if you can
- Plan out the use of the `co_*` keywords first in coro body
 - Then you can add the boilerplate
- The work goes in the coroutine body
- **Be very careful with object lifetime and ownership**
- You will have to write some sort of manager/queue
- Leave optimization for the very end



Be aware of object lifetime and ownership

- Be sure to return `suspend_always` from `final_suspend` if you want to get the result from `co_return`
- References to coroutine internals are valid for the life of the coroutine
- `Final_suspend` is final. You can't resume.



Conclusions

- Consider not using coroutines
 - Wait for standard library support in C++23 or C++26
 - Consider using a coroutine library if you can
- Coroutines are *not* subroutines – they're independent
- Plan out the use of the `co_*` keywords first in coro body
 - Then you can add the boilerplate
- The work goes in the coroutine body
- Be very careful with object lifetime and ownership
- You will have to write some sort of manager/queue
- Leave optimization for the very end

From Problem to Coroutine: Reducing I/O Latency

CHEINAN MARKS